

dns30ria5

September 14, 2023

```
[11]: # This Python 3 environment comes with many helpful analytics libraries
      ↪ installed
      # It is defined by the kaggle/python Docker image: https://github.com/kaggle/
      ↪ docker-python
      # For example, here's several helpful packages to load
      # Input data files are available in the read-only "../input/" directory
      # For example, running this (by clicking run or pressing Shift+Enter) will list
      ↪ all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

/kaggle/input/cardiovascular-disease-dataset/cardio_train.csv

1 Importing The Necessary Libraries

```
[12]: import pandas as pd
import numpy as np
import os, sys
import xgboost as xgb
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

2 Read CSV data into a DataFrame

```
[13]: df = pd.read_csv("/kaggle/input/cardiovascular-disease-dataset/cardio_train.
      ↪ csv", sep=';')
```

3 Checking the Shape of the data

```
[14]: df.shape
```

```
[14]: (70000, 13)
```

```
[29]: df.head(5)
```

```
[29]:
```

	id	age	gender	height	weight	ap_hi	ap_lo	cholesterol	gluc	smoke	\
0	0	18393	2	168	62.0	110	80	1	1	0	
1	1	20228	1	156	85.0	140	90	3	1	0	
2	2	18857	1	165	64.0	130	70	3	1	0	
3	3	17623	2	169	82.0	150	100	1	1	0	
4	4	17474	1	156	56.0	100	60	1	1	0	

	alco	active	cardio
0	0	1	0
1	0	1	1
2	0	0	1
3	0	1	1
4	0	0	0

4 Checking for any duplicate values

```
[16]: df.duplicated().sum()
```

```
[16]: 0
```

5 Checking for any Null Values

```
[17]: df.isnull().sum()
```

```
[17]: id          0
age          0
gender       0
height       0
weight       0
ap_hi        0
ap_lo        0
cholesterol  0
gluc         0
smoke        0
alco         0
active       0
cardio       0
```

dtype: int64

```
[18]: df.columns
```

```
[18]: Index(['id', 'age', 'gender', 'height', 'weight', 'ap_hi', 'ap_lo',  
        'cholesterol', 'gluc', 'smoke', 'alco', 'active', 'cardio'],  
        dtype='object')
```

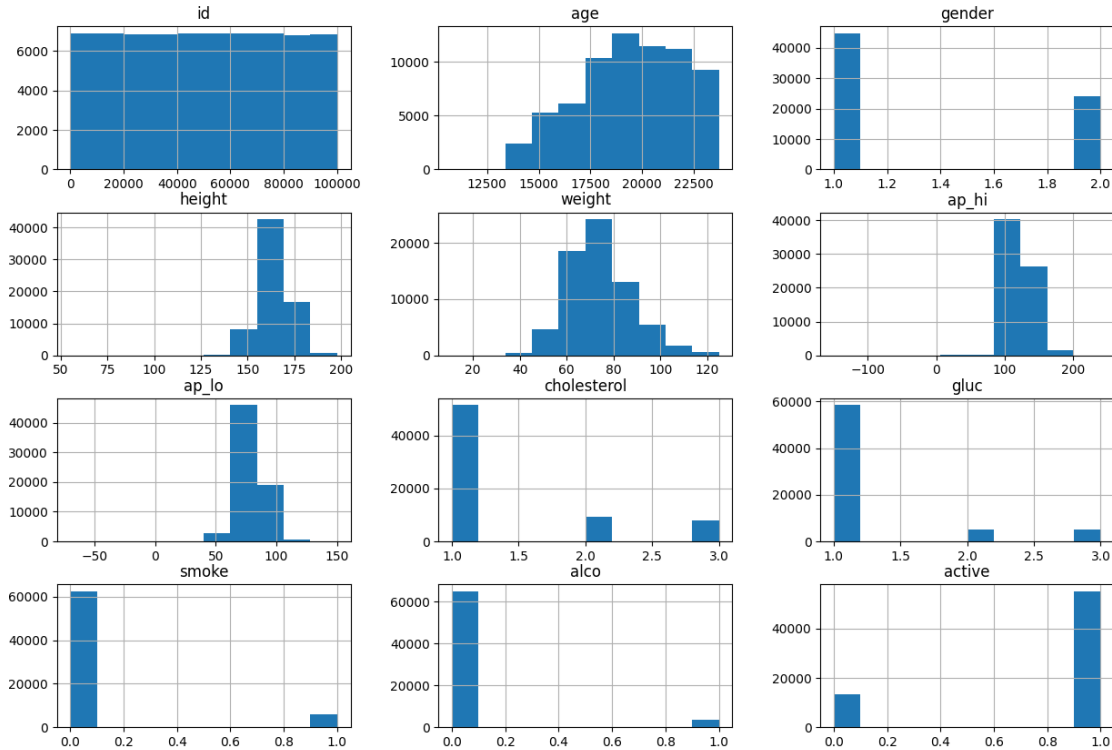
```
[19]: df.skew()
```

```
[19]: id            -0.001278  
age            -0.307055  
gender          0.630960  
height         -0.642187  
weight          1.012070  
ap_hi           85.296214  
ap_lo           32.114083  
cholesterol      1.587123  
gluc             2.397462  
smoke            2.905867  
alco             3.956607  
active          -1.529473  
cardio           0.001200  
dtype: float64
```

```
[20]: df=df[df.ap_hi<=250]  
df=df[df.ap_lo<=150]  
df=df[df.weight<=125]  
df=df[df.height<=200]
```

6 Plotting the Columns

```
[21]: x=df.loc[:,df.columns!='cardio'].values[:,1:]  
x1=df.loc[:,df.columns!='cardio']  
y=df.loc[:, 'cardio'].values  
y1=df.loc[:, 'cardio']  
x1.hist(figsize=(15,10))  
plt.show()
```



7 Scale the features to between -1 and 1

```
[22]: scaler=MinMaxScaler((-1,1))
      x1=scaler.fit_transform(x)
      y1=y
```

8 Split the dataset into Train and Test Datasets

```
[23]: xtrain,xtest,ytrain,ytest=train_test_split(x1, y1, test_size=0.
      ↪24,random_state=42)
```

9 Train the model using XGBoost Classifier

```
[24]: # Train the model
      from xgboost import XGBClassifier
      model=XGBClassifier()
      model.fit(xtrain,ytrain)
      predict=model.predict(xtest)
```


10 Confusion Matrix

```
[25]: from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

# Create the confusion matrix
cm = confusion_matrix(ytest, predict)

# Plot the confusion matrix heatmap
plt.figure(figsize=(6, 4))
fg = sns.heatmap(cm, annot=True, cmap="Reds", fmt='d')
figure = fg.get_figure()

# Add labels and title
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title("Output Confusion Matrix")
plt.show()
```



11 Printing the Confusion Matrix

```
[26]: print('True Positive Cases : {}'.format(cm[1][1]))  
      print('True Negative Cases : {}'.format(cm[0][0]))  
      print('False Positive Cases : {}'.format(cm[0][1]))  
      print('False Negative Cases : {}'.format(cm[1][0]))
```

```
True Positive Cases : 5520  
True Negative Cases : 6549  
False Positive Cases : 1797  
False Negative Cases : 2612
```

12 Calculating the Precision, Recall and F1-Score

```
[27]: pre = round(8486 / (8486 + 2767),3)  
      print("The Precision is:", pre)  
      rec = round(8486 / (8486 + 3794),3)  
      print("The Recall is:", rec)  
      f1_score = round(2 * (pre * rec) / (pre + rec),3)  
      print("The F1 Score is:", f1_score)
```

```
The Precision is: 0.754  
The Recall is: 0.691  
The F1 Score is: 0.721
```

13 Model Prediction Accuracy

```
[28]: print("The Model Accuracy is:",round(accuracy_score(ytest,predict)*100,3),"%")
```

```
The Model Accuracy is: 73.243 %
```